

Accuracy vs escalation rate — Speech Llama Questions (OpenAudioBench, n=250)  
probe v3, per-domain quantile thresholds; scored with OpenAudioBench's own judge (gpt-4o)

