

Score vs escalation rate — VoiceBench AlpacaEval (n=199)
probe v2, per-domain quantile thresholds; VoiceBench's own gpt-4o-mini judge prompt (1-5)

